

SpanDiff technical report

Summary

SpanDiff scores token-level semantic drift on Swiss federal pages in German, French, and Italian. It is a representation model, not a prompt. The submitted scorer is hidden state 15 of swiss-ai/Apertus-v1.5-8B. Subword states are mean-pooled onto whitespace tokens. The score is $\text{clamp}(1 - \max \text{cosine to the other side}, 0, 1)$. No tokens are generated. Apertus 70B was not run.

The official development script, with `--split dev`, printed Spearman de 0,296 ; fr 0,191 ; it 0,324. The unweighted mean of those three printed numbers is 0.270. Kendall's Tau-b printed de 0,238 ; fr 0,154 ; it 0,259. The stdout is `track_2a/docs/official_dev_stdout_L15.txt`.

Setup	de	fr	it	mean
Final layer, clamped DiffAlign	0.053	0.056	0.106	0.072
Hidden state 8 ablation	0.245	0.165	0.282	0.231
Hidden state 15, submitted	0.296	0.191	0.324	0.270

The mean column is that unweighted average, shown to 3 decimals. The rows use the printed coefficients. Decimal commas in the stdout are 0,296 and so on. Hidden state 8 is `docs/official_dev_stdout.txt`. The final layer is `docs/official_dev_stdout_last.txt`. All three grades used `--split dev` on 168 lines per language.

Method

Each side of a page is encoded alone. The model is swiss-ai/Apertus-v1.5-8B. Index 0 is the embedding. The submitted state is index 15, the residual stream after block 15. Index 32, after the final norm, is the final-layer baseline, not the submission.


A window holds at most 512 tokenizer tokens, including the leading beginning-of-sequence token. The development pass is 1619 such windows and 564,349 tokenizer tokens. Each of those windows has exactly one leading BOS. That token has an empty character offset and is not pooled. EOS, PAD, UNK, and multimodal special counts on that tokenizer-only census were 0. The census is `docs/special_tokens.txt`. It did not load weights.

Subwords are mean-pooled into the whitespace token they overlap most. An equal overlap stays on the earlier token. A whitespace token with no subword becomes a zero vector. Its cosine is defined as 0, so its score is 1. After the windows of one side are stitched, each token is scored against the other side's full sequence. The written score is $\text{clamp}(1 - \max \text{cosine}, 0, 1)$, rounded to 6 decimals.

The paper's DiffAlign takes the final hidden state and does not clamp, so a negative cosine can score above 1. SpanDiff clamps. The final-layer row in the table is the clamped version of that recipe on this 8B model. An unclamped rescore of the saved final-layer states was not produced. The paper also takes the maximum at subword level and then averages inside the word. SpanDiff pools the states first and then takes the maximum. The graded files keep that order.

Encoder lines contain only `id`, `text_a`, `text_b`, `labels_a`, and `labels_b`. Every label is in `[0, 1]`. The writer rejects `-1` and `NaN`. The label count equals the whitespace-token count. Generated tokens are 0, so no output filter is applied.

Layer curve



The figure plots the saved average-rank Spearman for every hidden state, separately for German, French, and Italian. The source table is docs/layer_map_by_lang_avg.txt. It uses the same token counts as the frozen-score tables (German 108,945, French 124,262, Italian 118,735) and drops gold labels of -1. It is not a second run of the official script. Do not treat a point on the curve as the printed contest score.

On that map, hidden state 8 reads 0.2495, 0.1662, and 0.2842. Hidden state 15 reads 0.2982, 0.1921, and 0.3241. German peaks at layer 14 (0.2983), a hair above layer 15. French peaks at layer 15. Italian peaks at layer 14 (0.3246), again a hair above layer 15. The final state, index 32, falls to 0.0484, 0.0564, and the German last block is the low end of that curve. A pooled ordinal Spearman across all three languages, in docs/layer_map_dev.txt, peaks at layer 14 (0.1408), with layer 15 at 0.1403, layer 8 at 0.1281, and the last block at -0.0358. That pooled number is one Spearman on 351,942 tokens. It is not the mean of three language scores.

The official prints sit next to this map, not on top of it. Hidden state 15 printed 0,296 ; 0,191 ; 0,324. Hidden state 8 printed 0,245 ; 0,165 ; 0,282. The final layer printed 0,053 ; 0,056 ; 0,106. The shape agrees: the middle of the network carries the signal, and the final state does not.

docs/token_heatmap.png shows one short development pair, admin_de_105, from the submitted layer-15 file. English and German tokens are colored by the clamped difference. It is an illustration, not a score.

Layer selection

The organizers publish a development train split of 134 documents per language and a development validation split of 34. Those files are gold_labels/dev/train and gold_labels/dev/val. Only the id field was read. The per-document Spearman's were then taken from docs/layer_map_dev.jsonl, which stores one ordinal Spearman per hidden state and does not store gold labels.

A document with constant gold has an undefined Spearman. Those documents were omitted: none of 134 German train pages, 1 of 34 German val, 2 of 134 French train, 1 of 34 French val, 1 of 134 Italian train, and 1 of 34 Italian val. The table is the unweighted mean of the remaining per-document Spearman's. It is not the official contest formula.

Slice	docs used	best layer	layer 8	layer 15	15 minus 8
de train	134/134	15	0.1623	0.1899	+0.0276
de val	33/34	15	0.1636	0.2003	+0.0367
fr train	132/134	15	0.0430	0.0798	+0.0367
fr val	33/34	15	0.0461	0.0801	+0.0340
it train	133/134	14	0.1509	0.1903	+0.0394
it val	33/34	15	0.1796	0.2066	+0.0270

Italian train is the exception. Its best layer is 14, at 0.1913, and layer 15 is 0.1903. An n-weighted mean of the same per-document values picks the same best layer in every row, including Italian train at layer 14. Layer 15 beats layer 8 by about 0.03 on every slice.

Layer 15 was then re-scored on all 168 development documents per language and graded with the official script. That grade, not the per-document map, is the submission. Hidden state 8 remains in predictions/ablation_l8.

Use of Apertus

One checkpoint: swiss-ai/Apertus-v1.5-8B. No second embedding model and no closed API. The development forward pass read 564,349 tokens and wrote 0 generated tokens. With the weights already cached, layer-15 scoring took 515 seconds on one Apple GPU. A cold load of the 8B weights needs about 16 GB. Apertus 70B was not run. It does not fit in 48 GB on this machine.

The text-only complement, 56 documents per language, was scored with the same layer-15 method so that a 224-document file can be assembled in the full-set id order. That forward pass was 207,075 tokens and 273 seconds. Those files were not graded. No test Spearman is claimed.

Data

Development inputs are the public SwissGov-RSD development pages, 168 documents per language. The 56-document complement and the 224-document id order are prediction files only. Gold labels for the sealed split are not in this repository. The SwissGov text stays CC-BY-4.0 and is not relicensed.

Evaluation

The summary table is the evaluation that should be quoted. It has three official development runs and no others. Layer-map numbers in the selection section answer why hidden state 15 was scored. They are not a substitute for the printed row.

Limitations

States are causal, so a token sees only its left context. An empty subword span scores 1, which inflates a few punctuation tokens. The final layer raises aligned tokens and different tokens together, so a last-layer DiffAlign score looks large and ranks poorly. Italian train, on the per-document map, prefers layer 14 by 0.001. The submitted choice is still layer 15 because validation prefers it in all three languages and the official full-development grade prefers it to layer 8. The unclamped paper formula was not rerun. 70B was not run. No test score is reported, and the test files must not be passed to the evaluator until the organizers say which split they score.

Reproducibility

From a checkout of this tree, make run-local grades predictions/dev with --split dev and does not load the model. make predict writes a new layer-15 development run under predictions/rerun. It needs the Apertus weights and a Hugging Face token only if those weights are not cached. The Docker image grades the same committed development files. It was built and graded under podman on linux/amd64. Docker is not installed on this Mac, so that image has not been run here.

Next steps

The open question is which file the organizers will score. Development, the 56-document complement, and the 224-document id order are all written. Only the development file has an official number. A 70B run is future work on a larger machine.

License

Report CC-BY-4.0. Code Apache-2.0. SwissGov text stays CC-BY-4.0, not relicensed.